

1. Preliminaries

Before starting on this assignment, please be sure to read the General Instructions that are on Canvas (under Files-->General Resources), and also make sure you are able to log in to the class PostgreSQL server. You'll get help on this in your Lab Section, not the Lectures, so *be sure to attend Lab Sections*.

2. Goal

The goal of the first assignment is to create a PostgreSQL data schema with 7 tables. That is all that is required in this assignment; don't do anything that's not required. The other Lab Assignments are much more difficult. In your Lab Sections, you may be given information about how to load data into a table and issue simple SQL queries, because that's fun, but loading data and issuing queries are **not** required in this assignment. (That will show up in the Lab2 assignment.)

3. Lab1 Description

3.1 Create PostgreSQL Schema Lab1

As we noted in the general instructions, you will create a Lab1 schema to set apart the database tables created in this lab from tables you will create in future labs, as well as from tables (and other objects) in the default (public) schema. Note that the meaning of schema here is specific to PostgreSQL, and distinct from the general meaning of schema. See [here](#) for more details on PostgreSQL schemas. You create the Lab1 schema using the following command:

```
CREATE SCHEMA Lab1;
```

[PostgreSQL makes all identifiers lowercase, unless you put them in quotation marks, e.g., "Lab1". But in CSE 180, you don't have to bother using quotation marks for identifiers. We use capitals in assignments for readability, but it's okay (and equivalent) if you use lab1 as the schema name.]

Now that you have created the schema, you want to make Lab1 be the default schema when you use psql. If you do not set Lab1 as the default schema, then you will have to qualify your table names with the schema name (e.g., by writing Lab1.customer, rather than just customer). To set the default schema, you modify your search path as follows. (For more details, see [here](#).)

```
ALTER ROLE username SET SEARCH_PATH to Lab1;
```

You will need to log out and log back in to the server for this default schema change to take effect. (Students **often forget** to do this, and then are surprised that their tables aren't in the expected schema.) To see your current SEARCH_PATH, enter:

```
SHOW SEARCH_PATH;
```

If you forget to do the ALTER ROLE for Lab1 to modify your search path then your schema will be the default schema, PUBLIC. Note that every student has their own database, so your PUBLIC schema won't be the same as any other student's PUBLIC schema, and the same is true for any other schema, such as Lab1.

3.2 Tables

You'll be creating tables for a very simplified version of a Newspaper Subscription database schema for the Daily Planet newspaper, with tables for SubscriptionKinds, Editions, Subscribers, Subscriptions, Holds, Articles, and ReadArticles. The data types and Referential Integrity for the attributes in these 7 Newspaper Subscription tables are described in the next section. No, this schema doesn't provide everything that a real-world Newspaper Subscription database would include, but it's a decent start.

Important: To receive full credit, you must use the attribute names as given, and the attributes must be in the order given. Also, the data types and referential integrity must match the specifications given in the next section. Follow directions; do not do more than you're asked to do in this assignment.

SubscriptionKinds(subscriptionMode, subscriptionInterval, rate, stillOffered)

Editions(editionDate, numArticles, numPages)

Subscribers(subscriberPhone, subscriberName, subscriberAddress)

Subscriptions(subscriberPhone, subscriptionStartDate, subscriptionMode, subscriptionInterval, paymentReceived)

Holds(subscriberPhone, subscriptionStartDate, holdStartDate, holdEndDate)

Articles(editionDate, articleNum, articleAuthor, articlePage)

ReadArticles(subscriberPhone, editionDate, articleNum, readInterval)

The underlined attribute (or attributes) identifies the Primary Key of each table. A table can only have one Primary Key, but that Primary Key may involve multiple attributes.

- A kind of subscription in SubscriptionKinds specifies the mode of the subscription (which can be 'D', 'P' or 'B', for digital, print or both) and the interval of that kind of subscription (which is a time interval, such as a week or a year), as well as the cost of that subscription type, and whether that kind of subscription is still being offered. For example, there could be a yearly print subscription whose cost is \$137.25 that is still being offered.
- An edition in Editions specifies the date of the edition, and the number of articles and pages in the edition.
- A subscriber in Subscribers specifies the phone, the name, and the address of that subscriber. subscriberPhone identifies the particular subscriber.
- A subscription in Subscriptions specifies the subscriber who has that subscription, the date the subscription started, the mode and interval of the subscription, and whether payment has been received for the subscription.
 - Any subscriberPhone that's in a Subscriptions row **must appear as** a subscriberPhone in the Subscribers table.
 - Any (subscriptionMode, subscriptionInterval) that's in a Subscriptions row **must appear as** a (subscriptionMode, subscriptionInterval) in the SubscriptionKinds table.

[Note that a subscriber could have more than one subscription.]

- A hold in Holds specifies the subscription being held (meaning that newspaper for that description shouldn't be delivered during the hold period). The subscription being held is identified by subscriberPhone and subscriptionStartDate. The start and end dates for the hold are also specified. (There might not be editions on those dates.)
 - Any (subscriberPhone, subscriptionStartDate) that's in a Holds row **must appear as** a (subscriberPhone, subscriptionStartDate) in the Subscriptions table.
- An article in Articles specifies the date of the edition in which the article appears, as well as its article number in that edition, its author, and the page on which it appears.
 - Any editionDate that's in an Articles row **must appear as** an editionDate in the Editions table.
- ReadArticles gives information about the time interval (readInterval) that a particular subscriber spent reading a particular article.
 - Any subscriberPhone that's in a ReadArticles row **must appear as** a subscriberPhone in the Subscribers table.
 - Any (editionDate, articleNum) that's in in a ReadArticles row **must appear as** a (editionDate, articleNum) in the Articles table.

In this assignment, you'll just have to create tables with the correct table names, attributes, data types, Primary Keys and Referential Integrity. "**Must appear as**" means that there's a Referential Integrity requirement. **Be sure not to forget Primary Keys and Referential Integrity when you do Lab1!**

3.2.1 Data types

Sometimes an attribute (such as state) appears in more than one table. Attributes that have the same attribute name might not have the same data type in all tables, but in our schema, they do.

- For subscriberPhone, numArticles, numPages, articleNum and articlePage use *integer*.
- For subscriptionMode, use *character with fixed length 1*.
- subscriberName and articleAuthor should be *character of variable length*, with maximum length 30.
- For subscriberAddress, use *character of variable length*, with maximum length 60.
- rate should be *numeric*, with at most 4 decimal digits to the left of the decimal point and 2 decimal digits after it.
- editionDate, subscriptionStartDate, holdStartDate and holdEndDate should be of type *date*.
- subscriptionInterval and readInterval should be of type *interval*.
- The stillOffered attribute, which indicates whether a kind of subscription is being offered to new subscribers, and the paymentReceived attribute, which indicates whether a payment for a subscription has been received, should be *boolean*.

You must write a CREATE TABLE statement for each of the 7 tables in Section 3.2. Write the statements in the same order that the tables are listed above. **Use the data types, Primary Keys and Referential Integrity described above.** You will lose credit if you do anything beyond that, even if you think that it's sensible. Save your statements in the file create_lab1.sql

PostgreSQL maps all SQL identifiers (e.g., table names and attributes) to lowercase. That's okay in your CSE 180 assignments. You won't lose points for Lab1 because Subscribers appears in the database as subscribers, and Articles appears as articles. It is possible to specify specific case choices for an identifier by putting that identifier inside double-quote symbols, e.g., as "Subscribers". But then every time you refer to that identifier, you'll have to use the double-quotes. "SUBSCRIBERS" is not the same identifier as "Subscribers", and neither of those is the same as Subscribers (written without double-quotes), which PostgreSQL maps to subscribers. We will use capitalization for readability, but we won't bother using double-quotes in our own Lectures and Lab Assignment or Exam solutions.

4. Testing

While you're working on your solution, it is a good idea to drop all objects from the schema every time you run the `create_lab1.sql` script, so you can start fresh. Dropping each object in a schema may be tedious, and sometimes there may be a particular order in which objects must be dropped. (Why?) The following command, which you should put at the top of your `create_lab1.sql` script, will drop your Lab1 schema (and all the objects within it), and then create the (empty) schema again:

```
DROP SCHEMA Lab1 CASCADE;  
CREATE SCHEMA Lab1;
```

The first statement will result in an error if the Lab1 Schema doesn't exist, but execution of your script will continue after that.

Before you submit your Lab1 solution, login to your database via `psql` and execute your `create_lab1.sql` script. As you'll learn in Lab Sections, the command to execute a script is: `\i <filename>`. Verify that every table has been created by using the command: `\d`. When you execute `\d`, the tables may be displayed in any order, not necessarily in the order in which you created them.

To verify that the attributes of each table are in the correct order, and that each attribute is assigned its correct data type use the following command: `\d <table>`.

We've supplied some load data that you can use to test your solution in the file `load_lab1.sql`. After you've created your tables, using the command: `\i create_lab1.sql`, you can load that data in `psql` by executing the command: `\i load_lab1.sql`. (Why will loading the data twice always result in errors?) If your solution fails on the load data, then it's likely that your solution has an error. But although testing can demonstrate that a program is buggy, testing cannot prove that a program is correct.

You do not have to develop your solution on `unix.ucsc.edu`, but please recognize that we'll run your solution on `unix.ucsc.edu`. If your solution fails on `unix.ucsc.edu`, you'll receive a poor grade, even if your solution worked in some other environment.

5. Submitting

1. Save your script as `create_lab1.sql`. You may add informative comments to your scripts if you want. Put any other information for the Graders in a separate README file that you may submit.
2. Zip the file(s) to a single file with name `Lab1_XXXXXXX.zip` where XXXXXXX is your 7-digit student ID. For example, if a student's ID is 1234567, then the file that this student submits for Lab1 should be named `Lab1_1234567.zip`.

If you have a README file (which is not required), you can use the Unix command:

```
zip Lab1_1234567 create_lab1.sql README
```

If you don't have a README file, to create the zip file you can use the Unix command:

```
zip Lab1_1234567 create_lab1.sql
```

(Of course, you should use **your own student ID**, not 1234567.) Submit a zip file, even if you only have one file.

Submit the zip file on Canvas under Assignment Lab1. Please be sure that you have access to Canvas for CSE 180. Registered students should automatically have access; students who are not registered in CSE 180 will not be submit solutions. You can replace your original solution, if you like, up to the Lab1 deadline. (Canvas will give the new file a slightly different name, but that's okay.) No students will be admitted to CSE 180 after the Lab1 due date.

If you are working on the UNIX timeshare and your zip file is located there, you will need to copy your file to your computer so that you can upload it to Canvas through your browser. To do that, you will need an FTP (File Transfer Protocol) client to securely transfer files from the UNIX timeshare. A widely used secure FTP client is Filezilla. Installation instructions are found in the site of [FileZilla](#) (make sure you install the distribution suitable for your operating system). After opening the Filezilla client, you will need to set the host field to `unix.ucsc.edu`, the username to your CruzId and the password to your Blue password, while the port number should be set to 22 (the default port for remote login). By clicking the Quickconnect button, if your credentials are correct, you will connect and be able to see the contents of your remote Unix folder at the right pane (under the title "Remote site"), while the left pane (under the title "Local site") will display the contents of your local file system. With the mouse, you can drag the file from the Unix folder and drop it to the desired location at your computer. This will transfer the file to your local machine, without erasing it from its original remote location. Filezilla is only one of several options for an FTP client. If you are finding it difficult to install the necessary tools and successfully do file transfers, you should promptly ask for help in the Lab Sections; **do not postpone this until the deadline date**. Computers in UCSC Unix Labs also have pre-installed SSH and FTP clients (PuTTY and PSFTP).

Other approaches to copy files includes using SCP (Secure Copy) and using Cut-and-Paste, where you copy the contents of the file from the unix system, and then paste contents into a file on your computer. Cut-and-Paste works well with for small files, but it's a hack that does not work well for large files.

The CSE 180 Teaching Assistants (Gavin Clark, Utkarsh Gupta and Dev Purandare) will discuss approaches to access unix remotely (SSH for Mac/Linux and PuTTY for Windows) and to move files to your computer (SCP for Mac/Linux and Filezilla for Windows/Mac/Linux) with you during Lab Sections. Attend your Lab Section to ensure that you know how to handle this correctly!

Lab1 is due by 11:59pm on Tuesday, October 17. **Late submissions will not be accepted; Canvas won't take them, nor will we.** Always check to make sure that your submission is on Canvas, and that you've submitted the correct file.

But as the Syllabus tells you, Lab Assignment grades don't count towards your CSE 180 Course Score! Even though Lab Assignments will be graded, the most important feedback students receive from Lab Assignment grades is about their errors. So be sure to submit your Lab Assignments on-time, so that you'll receive that feedback.

Yes, your CSE 180 Course Score won't be affected if you don't do any Lab Assignments, or if you do all your assignments using an LLM-based system such as ChatGPT. But those are very poor choices, since you won't be able to use LLM-based systems on CSE 180 Exams, or during job interviews.