

# Week 6 Tutoring

CSE 180





## Midterm 2

- Midterm 2 on Monday, 11/12
- Extra tutoring session on Thursday, 11/9, 2:00-3:00
- Includes Lecture 7 (and earlier) and Lecture 8 (except Triggers)
- Should be able to do everything on the practice midterm, but more will be included
  - Questions that were not on MT1: 5, 7, 9, 10, 11, 13, 16



# ACID

**Atomicity:** all or nothing; all changes in a transaction are performed or none

**Consistency:** domain of attribute must have certain restrictions and constraints

**Isolation:** whether or not multiple transactions trip over each other

**Durability:** after a transaction completes, changes to data persist



# Transactions

A transaction is a group of operations that should be executed atomically

The DBMS will execute each transaction in its entirety or not at all,  
“without transactions interfering with each other”.

START/BEGIN TRANSACTION <ISOLATION LEVEL >

<transaction>

COMMIT;



# Dirty Reads

**Dirty data** refers to data that is written by a transaction but has not yet been committed by the transaction.

A **dirty read** refers to the read of dirty data written by another transaction.

Consider the following transaction T that transfers an amount of money (\$X) from one account to another:

- 1) Add \$X to Account 2.
- 2) Test if Account 1 exists and has at least \$X in it.
  - a) If there is insufficient money, remove \$X from Account 2.
  - b) Otherwise, subtract \$X from Account 1.

- Transaction T1: Transfers \$150 from A1 to A2.
- Transaction T2: Transfers \$250 from A2 to A3.
- Initial values: A1: \$100, A2: \$200, A3: \$300.

- 1) A1: \$100 , A2: \$200, A3: \$300
- 2) A1: \$100 , A2: \$200, A3: \$300
- 3) A1: \$100 , A2: \$350, A3: \$300
- 4) A1: \$100 , A2: \$350, A3: \$550
- 5) A1: \$100 , A2: \$100, A3: \$550



# Isolation Levels

## READ UNCOMMITTED

- Allows dirty reads

## READ COMMITTED (default)

- Only clean reads
- May read data committed by different transactions (no repeatable reads)

## REPEATABLE READ (default)

- Repeated queries will retrieve the same value even if it changed by another transaction



# Isolation Levels Cont.

## SERIALIZABLE

- As if one-by-one
- Preserves consistency
- Worse response time and throughput

## SNAPSHOT ISOLATION

- Along with read committed and repeatable read, them most common
- Better performance but worse consistency than SERIALIZABLE
- Transaction read data as it existed when transaction began



# Isolation Levels

Isolation Level	Clean Reads	Repeatable Reads	Simultaneous Existence
READ UNCOMMITTED	N	N	N
READ COMMITTED	Y	N	N
REPEATABLE READ	Y	Y	N
SERIALIZABLE	Y	Y	Y





# Snapshot Isolation

- Transaction reads data as it existed when transaction began (hence reads are repeatable).
- Prevents Write conflicts
  - ... but not on Read/Write conflicts between transactions (which Serializability DOES do)



# SQL Languages



# SQL Languages

## Data Manipulation Language (DML)

- Access and modify data
- SELECT, INSERT, DELETE, UPDATE

## Data Definition Language (DDL)

- Modify structure of data
- CREATE, DROP, ALTER

## Data Control Language (DCL)

- Control access to the data (security)
- GRANT, REVOKE



# CREATE/DROP TABLE

```
CREATE TABLE <table name> (
```

```
    Attributes ...
```

```
)
```

```
DROP TABLE <table name>;
```



# ALTER TABLE

- Adding a column to a table:

ALTER TABLE <table name> ADD <attribute name> <attribute type> <DEFAULT/NOT NULL>;

Dropping a column from a table:

ALTER TABLE <table name> DROP <column name>;

- In some SQL systems, dropping a column isn't allowed.



# VIEWS

- Views help with logical data independence, allowing you to retrieve it if it matches the description in the view
- Advantages:
  - short-hand/encapsulation
  - Re-use
  - Authorization
  - Logical data independence

CREATE VIEW <view name>(<attribute names>) AS

SELECT <attributes>

FROM R...

WHERE <conditions>



# INDEX

- We use indexes so we can easily search tables
  - Can also make INSERT, DELETE, UPDATE slower
- If a table is updated, all indexes are immediately updated within the transaction
- Preserves physical data independence
  - SQL statements can be executed regardless of which indexes (if any) exist in the database
- What gets impacted when indexes are created or dropped?
  - Performance of SQL statements
  - Some may run faster, some may run slower



# Constraints and Triggers

\* Triggers will not be included on Midterm 2







# Kinds of Constraints

- Primary Key/Unique constraints
- Foreign Key, or referential-integrity constraints
- Attribute-based constraints
  - Constrain values of a particular attribute
- Tuple-based constraints
  - Relationship among components of tuple
- Assertions
  - Any SQL boolean expression (not implemented in most relational DBMS, not discussed in this lecture)



# Adding/Dropping FK Constraints

ALTER TABLE R

ADD CONSTRAINT <foreign key name> FOREIGN KEY (<attributes>)

REFERENCES R(<attributes>)

ALTER TABLE R

DROP CONSTRAINT <foreign key name>



# INSERT/UPDATE with FK Constraints

Options:

1. RESTRICT: Reject the modification; the default
2. CASCADE: Makes the same changes in all tables that are referenced
3. SET NULL: Change the corresponding tuples that are referenced to NULL



# CHECK Constraints

```
ALTER TABLE R
```

```
ADD CONSTRAINT <constraint name>
```

```
CHECK (<condition>);
```